Week 2 Python Assessment

This Jupyter Notebook is auxillary to the following assessment in this week. To complete this assessment, you will complete the 7 questions outlined in this document and use the output from your python cells as answers.

Your goal of this assignment is to construct regression and logistics models and interpret model paramters.

Run the following cell to initialize your environment and begin the assessment.

In [1]:
```
#### RUN THIS

import warnings
warnings.filterwarnings('ignore')

import numpy as np
import statsmodels.api as sm
import pandas as pd

from sklearn.datasets import load_boston
boston_dataset = load_boston()

boston = pd.DataFrame(data=boston_dataset.data, columns=boston_dataset.feature_names)
boston["MEDV"] = boston_dataset.target

url = "nhanes_2015_2016.csv"
NHANES = pd.read_csv(url)
vars = ["BPXSY1", "RIDAGEYR", "RIAGENDR", "RIDRETH1", "DMDEDUC2", "BMXBMI", "SMQ020"]
NHANES = NHANES[vars].dropna()
NHANES["smq"] = NHANES.SMQ020.replace({2: 0, 7: np.nan, 9: np.nan})
NHANES["RIAGENDRx"] = NHANES.RIAGENDR.replace({1: "Male", 2: "Female"})
NHANES["DMDEDUC2x"] = NHANES.DMDEDUC2.replace({1: "lt9", 2: "x9_11", 3: "HS", 4: "SomeCollege",5: "College", 7: np.nan, 9: np.nan}

np.random.seed(123)
```

Now that your notebook is ready, begin answering the questions below.

Questions 1-3

The first three questions will be utilizing the Boston housing dataset seen in week 1.

Here is the description for each column:

- **CRIM:** Per capita crime rate by town
- **ZN:** Proportion of residential land zoned for lots over 25,000 sq. ft
- **INDUS:** Proportion of non-retail business acres per town
- **CHAS:** Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- **NOX:** Nitric oxide concentration (parts per 10 million)
- **RM:** Average number of rooms per dwelling
- **AGE:** Proportion of owner-occupied units built prior to 1940
- **DIS:** Weighted distances to five Boston employment centers
- **RAD:** Index of accessibility to radial highways
- **TAX:** Full-value property tax rate per $\$10,000$
- **PTRATIO:** Pupil-teacher ratio by town
- **B:** $1000(Bk — 0.63)^2$, where Bk is the proportion of [people of African American descent] by town
- **LSTAT:** Percentage of lower status of the population
- **MEDV:** Median value of owner-occupied homes in $\$1000$s

Uncomment and run the following code to generate a simple linear regression and output the model summary:

In [2]:
```
model = sm.OLS.from_formula("MEDV ~ RM + CRIM", data=boston)
result = model.fit()
result.summary()
```

Out[2]:

OLS Regression Results

Dep. Variable:	MEDV	**R-squared:**	0.541
Model:	OLS	**Adj. R-squared:**	0.539
Method:	Least Squares	**F-statistic:**	295.9
Date:	Wed, 14 Apr 2021	**Prob (F-statistic):**	1.15e-85
Time:	03:44:03	**Log-Likelihood:**	-1643.5
No. Observations:	506	**AIC:**	3293.
Df Residuals:	503	**BIC:**	3306.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	-29.3017	2.592	-11.303	0.000	-34.395	-24.208
RM	8.3975	0.406	20.706	0.000	7.601	9.194
CRIM	-0.2618	0.033	-7.899	0.000	-0.327	-0.197

Omnibus:	170.471	**Durbin-Watson:**	0.805
Prob(Omnibus):	0.000	**Jarque-Bera (JB):**	1034.461
Skew:	1.331	**Prob(JB):**	2.34e-225
Kurtosis:	9.479	**Cond. No.**	92.2

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Utilizing the above output, answer the following three questions:

Question 1 (You'll answer this question within the quiz that follows this notebook)

What is the value of the coefficient for predictor **RM**?

Question 2 (You'll answer this question within the quiz that follows this notebook)

Are the predictors for this model statistically significant, yes or no? (Hint: What are their p-values?)

Run the following code for question 3:

In [3]:
```
## For Question 3
model = sm.OLS.from_formula("MEDV ~ RM + CRIM + LSTAT", data=boston)
result = model.fit()
result.summary()
```

Out[3]:

OLS Regression Results

Dep. Variable:	MEDV	**R-squared:**	0.646
Model:	OLS	**Adj. R-squared:**	0.644
Method:	Least Squares	**F-statistic:**	304.9
Date:	Wed, 14 Apr 2021	**Prob (F-statistic):**	1.19e-112
Time:	03:44:16	**Log-Likelihood:**	-1577.8
No. Observations:	506	**AIC:**	3164.
Df Residuals:	502	**BIC:**	3180.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	-2.4978	3.165	-0.789	0.430	-8.717	3.721
RM	5.2092	0.442	11.785	0.000	4.341	6.078
CRIM	-0.1011	0.032	-3.162	0.002	-0.164	-0.038
LSTAT	-0.5804	0.048	-12.201	0.000	-0.674	-0.487

Omnibus:	171.189	**Durbin-Watson:**	0.822
Prob(Omnibus):	0.000	**Jarque-Bera (JB):**	623.248
Skew:	1.531	**Prob(JB):**	4.61e-136
Kurtosis:	7.492	**Cond. No.**	216.

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Question 3 (You'll answer this question within the quiz that follows this notebook)

What happened to our R-Squared value when we added the third predictor **LSTAT** to our initial model?

Question 4 (You'll answer this question within the quiz that follows this notebook)

What type of model should we use when our target outcome, or dependent variable is continuous?

Questions 5-6

The next two questions will involve the NHANES dataset.

Uncomment and run the following code to generate a logistics regression and output the model summary:

In [4]:
```
model = sm.GLM.from_formula("smq ~ RIAGENDRx + RIDAGEYR + DMDEDUC2x", family=sm.families.Binomial(), data=NHANES)
result = model.fit()
result.summary()
```

Out[4]:

Generalized Linear Model Regression Results

Dep. Variable:	smq	**No. Observations:**	5093
Model:	GLM	**Df Residuals:**	5086
Model Family:	Binomial	**Df Model:**	6
Link Function:	logit	**Scale:**	1.0000
Method:	IRLS	**Log-Likelihood:**	-3201.2
Date:	Wed, 14 Apr 2021	**Deviance:**	6402.4
Time:	03:44:29	**Pearson chi2:**	5.10e+03
No. Iterations:	4	**Covariance Type:**	nonrobust

	coef	std err	z	P>|z|	[0.025	0.975]
Intercept	-2.3060	0.114	-20.174	0.000	-2.530	-2.082
RIAGENDRx[T.Male]	0.9096	0.060	15.118	0.000	0.792	1.028
DMDEDUC2x[T.HS]	0.9434	0.090	10.521	0.000	0.768	1.119
DMDEDUC2x[T.SomeCollege]	0.8322	0.084	9.865	0.000	0.667	0.998
DMDEDUC2x[T.t9]	0.2662	0.109	2.438	0.015	0.052	0.480
DMDEDUC2x[T.x9_11]	1.0986	0.107	10.296	0.000	0.889	1.308
RIDAGEYR	0.0183	0.002	10.582	0.000	0.015	0.022

Question 5 (You'll answer this question within the quiz that follows this notebook)

Which of our predictors has the largest coefficient?

Question 6 (You'll answer this question within the quiz that follows this notebook)

Which values for DMDEDUC2x and RIAGENDRx are represented in our intercept, or what is our reference level?

Question 7 (You'll answer this question within the quiz that follows this notebook)

What model should we use when our target outcome, or dependent variable is binary, or only has two outputs, 0 and 1.